


User writes normally. LLM generates ALL responses.
KIM only provides context and validation.



Interaction Flow

1. User writes message to their LLM
2. LLM calls `get_context(query)`
KIM returns profile + relevant examples
3. LLM generates draft response
(using KIM's context)
4. LLM calls `check_draft(draft)`
KIM validates against profile
store for future retrieval
5. LLM revises if needed
6. LLM calls `log_output(final)`
KIM stores for future retrieval

Key insight:
The LLM thinks `check_draft` is an independent external service. It gets unbiased feedback because KIM validates with fresh eyes (profile rules only, no generated context). (Claude and advanced clients only)

- If the LLM supports MCP sampling, KIM can orchestrate multi-step flows:
- Fresh context per “agent” role
 - KIM controls the loop
 - Better isolation between steps
 - Same user LLM does the compute

Falls back to Layer 1 (Anleitung) for simpler MCP clients.